



OpenCL介绍



什么是OpenCL

OpenCL (Open Computing Language) , 即开放运算语言, 是一个统一的开放式的开发平台。

OpenCL是首个提出的并行开发的开放式的、兼容的、免费的标准。

它的目的是为异构系统通用提供统一开发平台。

OpenCL最初是由苹果公司设想和开发, 并在与AMD, IBM, 英特尔和NVIDIA技术团队的合作之下初步完善。随后, 苹果将这一草案提交至Khronos Group。

参见

<https://www.jianshu.com/p/b55552ee61ac>

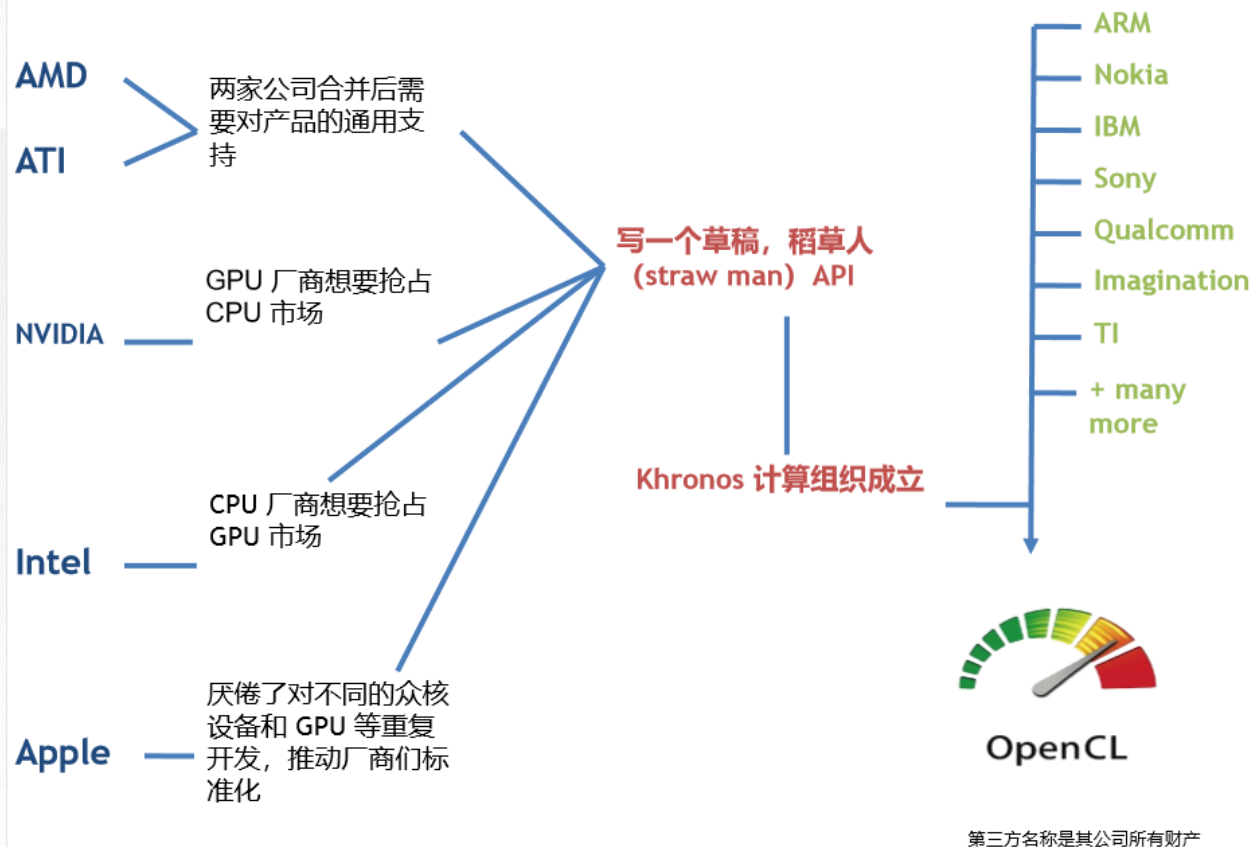
<https://www.cnblogs.com/wangshide/archive/2012/01/07/2315830.html>

<https://blog.csdn.net/eric41050808/article/details/10210025>



OpenCL

OpenCL 的起源



现代的计算平台包括

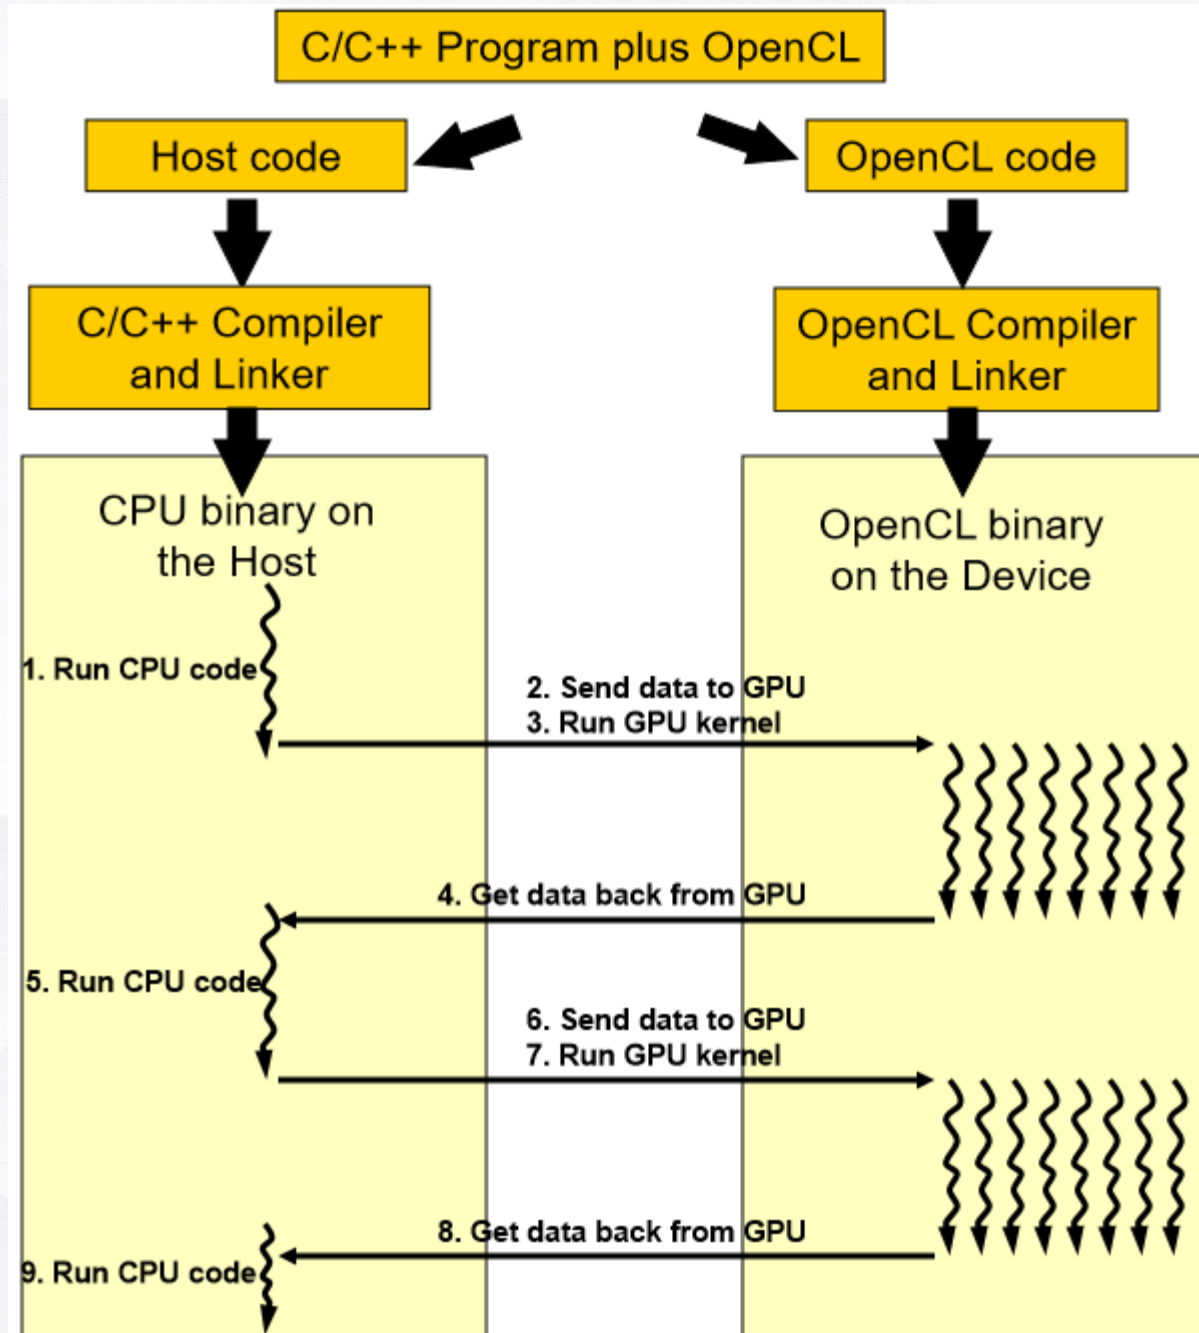
- 一个或多个 CPU
- 一个或多个 GPU
- DSP 处理器
- 加速器(accelerator)
- ... 其他的

OpenCL 允许开发者写一份可以移植的代码, 在不同平台上都可以使用所有的资源。



OpenCL

OpenCL的例子





什么是CUDA

CUDA (Compute Unified Device Architecture) , 即计算统一设备架构

显卡厂商NVIDIA推出的运算平台, 也是一种由NVIDIA推出的通用并行计算架构

用于NVIDIA设备, 能够解决复杂的计算问题

参见

<https://www.cnblogs.com/skyfsm/p/9673960.html>



与CUDA的对比

	CUDA	OpenCL
是什么	硬件架构，指令集架构（ISA），编程语言，API，SDK，工具	开放的API和语言标准
是否自由	非自由	自由，买断式授权
引入时间	2006	2008
SDK供应商	NVIDIA	具体根据企业
实现企业	NVIDIA	Apple、NVIDIA、AMD、IBM
支持系统	Windows, Linux, Mac OS X; 32 and 64-bit	依赖具体企业
支持设备类型	NVIDIA GPU	多种类型
支持嵌入式设备	不支持	支持



OpenCL

CPU

```
// function to add the elements of two arrays
void add(int n, double *x, double *y, double *z)
{
    for (int i = 0; i < n; i++)
        z[i] = x[i] + y[i];
}
```

OpenCL

```
__kernel void adder(__global const float* a, __global const float* b,
__global float* result)
{
    int idx = get_global_id(0);
    result[idx] = a[idx] + b[idx];
}
```

CUDA

```
__global__
void add(int n, double *x, double *y, double *z)
{
    // compute using the whole grid at once
    // the reuse the threads in the same grid
    // assume we have 1k elements to compute, and a stride of 100
    // thread 0 is responsible for c[0] = a[0] + b[0]
    //                                     c[100] = a[100] + b[100]...
    // thread 1 is responsible for c[1] = a[1] + b[1]
```


```

//                                     c[101] = a[101] + b[101]...
// and so on
int index = blockIdx.x * blockDim.x + threadIdx.x;
int stride = blockDim.x * gridDim.x;
for(int i=index; i<n; i+=stride)
    z[i] = x[i] + y[i];
}
```



OpenCL

与CUDA的一些对比





OpenCL

CUDA

```
int main(void)
{
    int N = 1<<20;
    double *x, *y,*z;

    // Allocate Unified Memory - accessible from CPU or GPU
    cudaMallocManaged(&x, N*sizeof(double));
    cudaMallocManaged(&y, N*sizeof(double));
    cudaMallocManaged(&z, N*sizeof(double));

    // initialize x and y arrays on the host
    for (int i = 0; i < N; i++) {
        x[i] = std::rand();
        y[i] = std::rand();
    }

    // Run kernel on 1M elements on the GPU
    int blockSize = 256;
    int numBlocks = (N+blockSize-1)/blockSize;
    add<<<numBlocks, blockSize>>>(N, x, y, z);

    // Wait for GPU to finish before accessing on host
    cudaDeviceSynchronize();

    // Check for errors
    double maxError = 0;
    for (int i = 0; i < N; i++)
        maxError = fmax(maxError, fabs(z[i]-x[i]-y[i]));
    std::cout << "Max error: " << maxError << std::endl;

    // Free memory
    cudaFree(x);
    cudaFree(y);
    cudaFree(z);

    return 0;
}
```



OpenCL

1. 引入头文件

```
1 #ifdef __APPLE__
2 #include <OpenCL/opencl.h>
3 #else
4 #include <CL/cl.h>
5 #endif
```

2. 获取系统上所有OpenCL平台

```
1 cl_int err;
2 cl_uint num;
3 err = clGetPlatformIDs(0, 0, &num);
4 if(err != CL_SUCCESS) {
5     std::cerr << "Unable to get platforms\n";
6     return 0;
7 }
```

3. 取得平台的ID, 用于建立OpenCL上下文 (context)

```
1 std::vector<cl_platform_id> platforms(num);
2 err = clGetPlatformIDs(num, &platforms[0], &num);
3 if(err != CL_SUCCESS) {
4     std::cerr << "Unable to get platform ID\n";
5     return 0;
6 }
```

4. 建立OpenCL上下文

```
1 cl_context_properties prop[] = { CL_CONTEXT_PLATFORM,
2     reinterpret_cast<cl_context_properties>(platforms[0]), 0 };
3 cl_context context = clCreateContextFromType(prop, CL_DEVICE_TYPE_DEFAULT,
4     NULL, NULL, NULL);
5 if(context == 0) {
6     std::cerr << "Can't create OpenCL context\n";
7     return 0;
8 }
```

5. 取得OpenCL设备列表

```
1 size_t cb;
2 clGetContextInfo(context, CL_CONTEXT_DEVICES, 0, NULL, &cb);
3 std::vector<cl_device_id> devices(cb / sizeof(cl_device_id));
4 clGetContextInfo(context, CL_CONTEXT_DEVICES, cb, &devices[0], 0);
```

6. 选定OpenCL设备

```
1 clGetDeviceInfo(devices[0], CL_DEVICE_NAME, 0, NULL, &cb);
2 std::string devname;
3 devname.resize(cb);
4 clGetDeviceInfo(devices[0], CL_DEVICE_NAME, cb, &devname[0], 0);
5 std::cout << "Device: " << devname.c_str() << "\n";
```

7. 建立命令队列 (接收操作)

```
1 cl_command_queue queue = clCreateCommandQueue(context, devices[0], 0, 0);
2 if(queue == 0) {
3     std::cerr << "Can't create command queue\n";
4     clReleaseContext(context);
5     return 0;
6 }
```




OpenCL

8. 配置内存并复制数据

```
1  const int DATA_SIZE = 1048576;
2  std::vector<float> a(DATA_SIZE), b(DATA_SIZE), res(DATA_SIZE);
3  for(int i = 0; i < DATA_SIZE; i++) {
4      a[i] = std::rand();
5      b[i] = std::rand();
6  }
7
8  cl_mem cl_a = clCreateBuffer(context, CL_MEM_READ_ONLY |
9  CL_MEM_COPY_HOST_PTR, sizeof(cl_float) * DATA_SIZE, &a[0], NULL);
10 cl_mem cl_b = clCreateBuffer(context, CL_MEM_READ_ONLY |
11 CL_MEM_COPY_HOST_PTR, sizeof(cl_float) * DATA_SIZE, &b[0], NULL);
12 cl_mem cl_res = clCreateBuffer(context, CL_MEM_WRITE_ONLY, sizeof(cl_float)
13 * DATA_SIZE, NULL, NULL);
14 if(cl_a == 0 || cl_b == 0 || cl_res == 0) {
15     std::cerr << "Can't create OpenCL buffer\n";
16     clReleaseMemObject(cl_a);
17     clReleaseMemObject(cl_b);
18     clReleaseMemObject(cl_res);
19     clReleaseCommandQueue(queue);
20     clReleaseContext(context);
21     return 0;
22 }
```

创建program object

```
1  cl_program program = load_program(context, "shader.cl");
2  if(program == 0) {
3      std::cerr << "Can't load or build program\n";
4      clReleaseMemObject(cl_a);
5      clReleaseMemObject(cl_b);
6      clReleaseMemObject(cl_res);
7      clReleaseCommandQueue(queue);
8      clReleaseContext(context);
9      return 0;
10 }
```

取得程序中函数的进入点

```
1  cl_kernel adder = clCreateKernel(program, "adder", 0);
2  if(adder == 0) {
3      std::cerr << "Can't load kernel\n";
4      clReleaseProgram(program);
5      clReleaseMemObject(cl_a);
6      clReleaseMemObject(cl_b);
7      clReleaseMemObject(cl_res);
8      clReleaseCommandQueue(queue);
9      clReleaseContext(context);
10     return 0;
11 }
```

10. 执行OpenCL kernel

设定参数

```
1  clSetKernelArg(adder, 0, sizeof(cl_mem), &cl_a);
2  clSetKernelArg(adder, 1, sizeof(cl_mem), &cl_b);
3  clSetKernelArg(adder, 2, sizeof(cl_mem), &cl_res);
```

开始执行

```
1  size_t work_size = DATA_SIZE;
2  err = clEnqueueNDRangeKernel(queue, adder, 1, 0, &work_size, 0, 0, 0, 0);
3  if(err == CL_SUCCESS) {
4      err = clEnqueueReadBuffer(queue, cl_res, CL_TRUE, 0, sizeof(float) *
5      DATA_SIZE, &res[0], 0, 0, 0);
6  }
```



与CUDA的一些对比

CUDA

gridDim

blockIdx

blockDim

gridDim * blockDim

threadIdx

blockIdx * blockDim + threadIdx

OpenCL

get_num_groups()

get_group_id()

get_local_size()

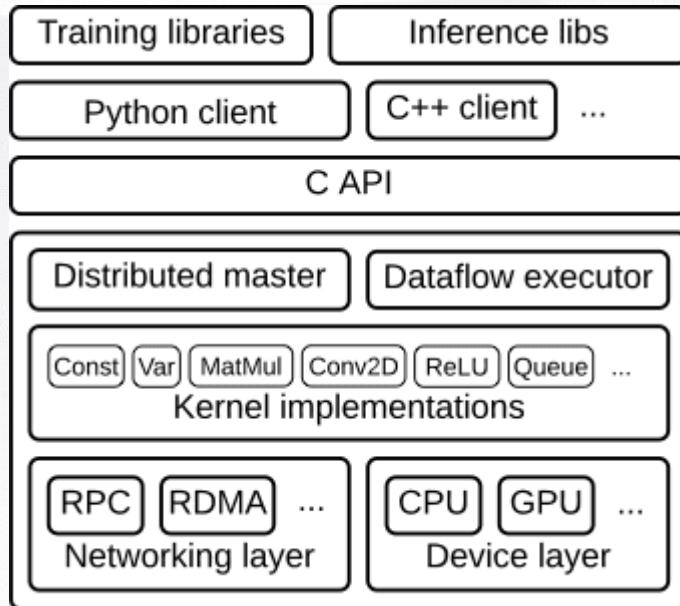
get_global_size()

get_local_id()

get_global_id()



Tensorflow kernel



kernel是Tensorflow中操作的具体实现

例如图中的Matmul（矩阵乘法）、Conv2D（二维卷积）

这些操作可以在CPU上进行，也可以在GPU上运行

<https://blog.csdn.net/u013510838/article/details/84103503>



相关资料

CUDA与OpenCL

https://www.sharcnet.ca/help/index.php/Porting_CUDA_to_OpenCL (API对比)

<https://www.cnblogs.com/huliangwen/p/5003504.html> (概念对比)

<https://blog.csdn.net/ijuliet/article/details/4631214> (代码对比)

OpenCL代码详情

<https://blog.csdn.net/breakawayroad/article/details/8227450> (修饰符)

<https://www.khronos.org/registry/OpenCL/sdk/1.2/docs/man/xhtmll> (API)



*Thank you for
watching*

